

Università degli Studi Roma Tre Dipartimento di Ingegneria Computer Networks Research Group

kathara lab

rip

Version	1.0
Author(s)	G. Di Battista, M. Patrignani, M. Pizzonia, F. Ricci, M. Rimondini
E-mail	contact@kathara.org
Web	http://www.kathara.org/
Description	experiences with the ripv2 distance vector routing protocol – kathara version of the rip lab of netkit (vers. 2.4)

copyright notice

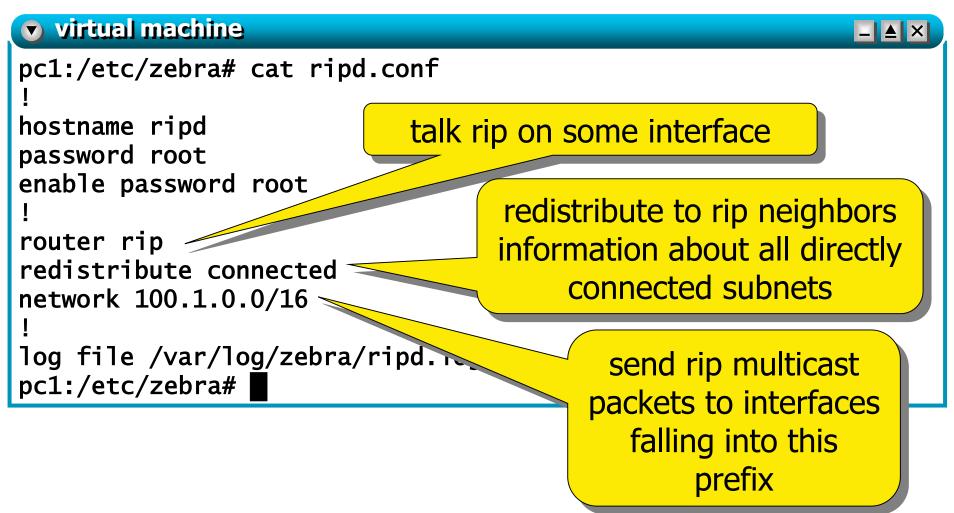
- All the pages/slides in this presentation, including but not limited to, images, photos, animations, videos, sounds, music, and text (hereby referred to as "material") are protected by copyright.
- This material, with the exception of some multimedia elements licensed by other organizations, is property of the authors and/or organizations appearing in the first slide.
- This material, or its parts, can be reproduced and used for didactical purposes within universities and schools, provided that this happens for non-profit purposes.
- Information contained in this material cannot be used within network design projects or other products of any kind.
- Any other use is prohibited, unless explicitly authorized by the authors on the basis of an explicit agreement.
- The authors assume no responsibility about this material and provide this material "as is", with no implicit or explicit warranty about the correctness and completeness of its contents, which may be subject to changes.
- This copyright notice must always be redistributed together with the material, or its portions.

last update: Sept 2018

routing protocols

- routing protocols are used to automatically update the routing tables
- they fall into two main cathegories:
 - link-state routing protocols
 - approach: send the minimum information to everyone
 - each router reconstructs the whole network graph and computes a shortest path tree to all destinations
 - examples: is-is, ospf
 - distance-vector routing protocols
 - approach: send all your information to a few
 - update your routing information based on what you hear
 - examples: rip, bgp
- in this lab we will see an example of RIPv2 protocol on quagga boxes

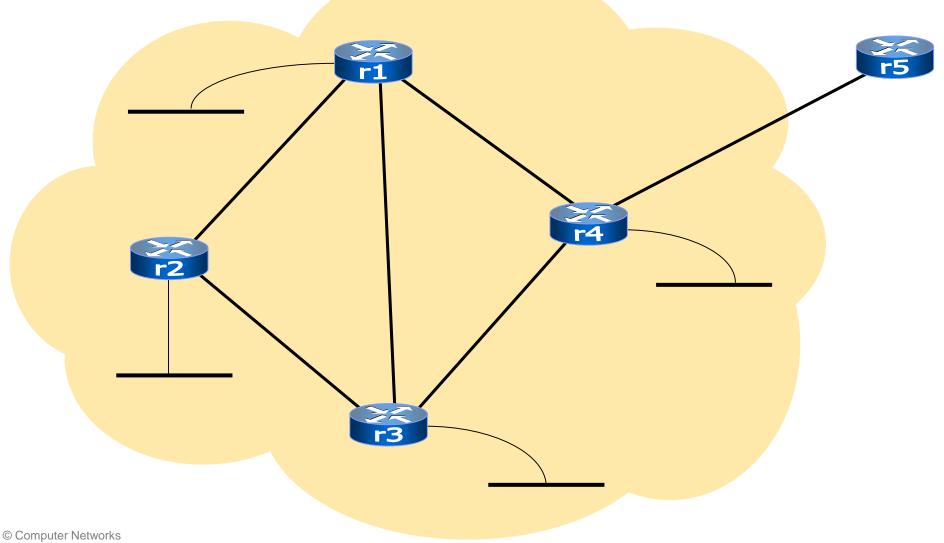
sample ripd configuration file (ripd.conf)



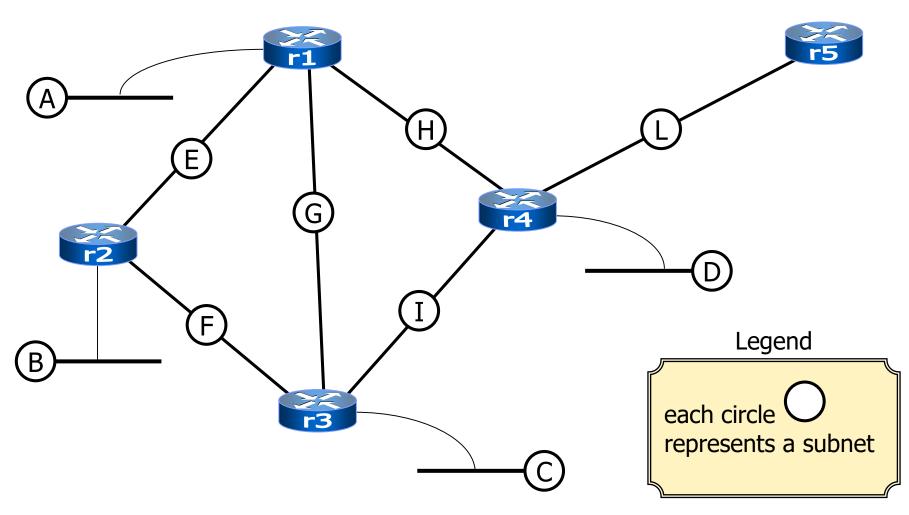
about redistribute connected

- by default (i.e., without further configuration) RIP already propagates information about directly connected subnets...
 - ...attached to RIP-speaking interfaces only
- redistribute connected forces RIP to propagate information about all connected subnets
- the semantic of redistribute connected applies to all routing protocols
- the default behavior does not: some protocols (e.g., bgp) are lazier, and do not propagate anything unless explicitly told to do so

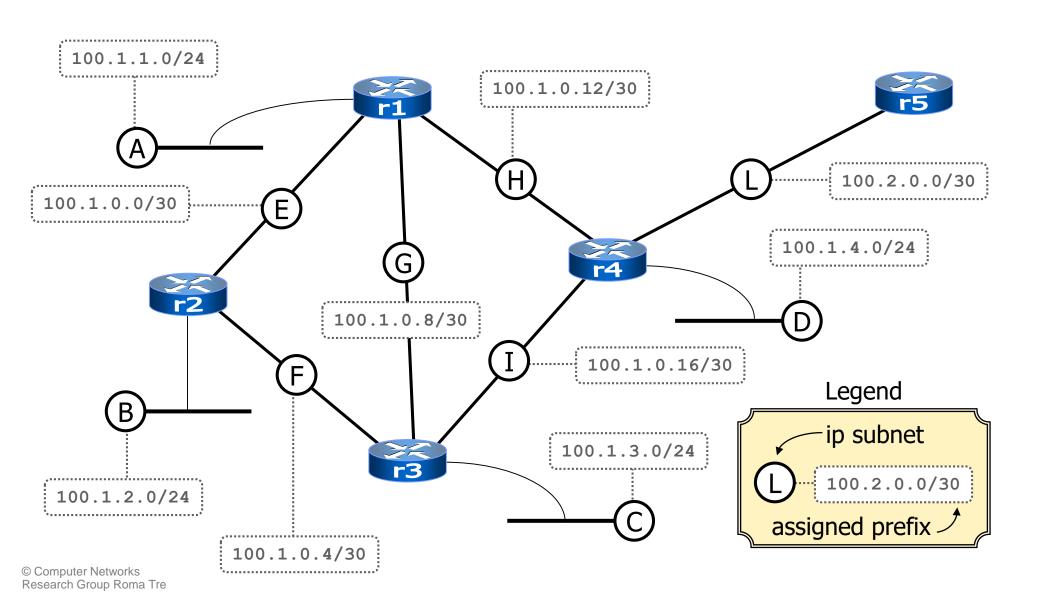
a small network connected to the Internet



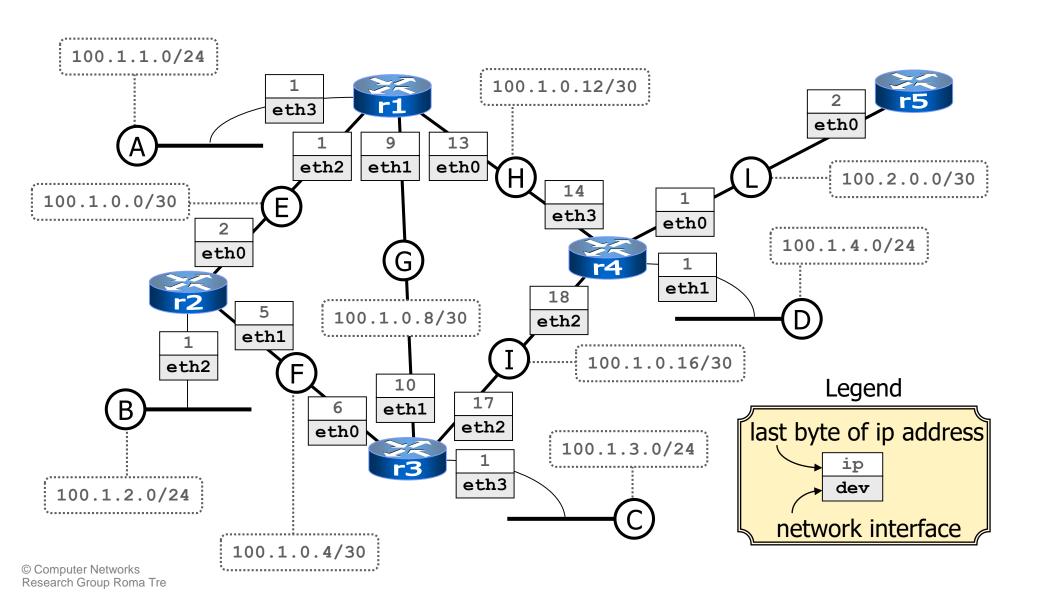
the involved ip subnets



assigning ip numbers to subnets



assigning ip numbers to interfaces



launching the lab script



- the lab configuration is such that
 - five virtual hosts are created and connected to the right collision domains (virtual hubs)
 - for each virtual host
 - network interfaces are automatically configured
 - configuration files /etc/quagga/daemons, /etc/quagga/zebra.conf, and /etc/quagga/ripd.conf are updated
 - the zebra routing daemon is <u>not</u> automatically started

checking connectivity

towards a directly connected destination

```
r4:~# ping 100.1.0.13
PING 100.1.0.13 (100.1.0.13) 56(84) bytes of data.
64 bytes from 100.1.0.13: icmp_seq=1 ttl=64 time=1.23 ms
64 bytes from 100.1.0.13: icmp_seq=2 ttl=64 time=0.592 ms
64 bytes from 100.1.0.13: icmp_seq=3 ttl=64 time=0.393 ms

--- 100.1.0.13 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2032ms
rtt min/avg/max/mdev = 0.393/0.741/1.238/0.360 ms
r4:~#
```

checking connectivity

towards a remote destination

```
r4:~# ping 100.1.2.1
connect: Network is unreachable
r4:~# ■
```

what's going on?

examining the kernel routing table

```
▼ r4
                                                                         _ ≜ ×
r4:~# route
Kernel IP routing table
Destination
                                                Flags Metric Ref
                                                                    Use Iface
                                Genmask
                Gateway
100.1.0.16
                                255.255.255.252 U
                                                                       0 eth2
                                255.255.255.252 U
100.2.0.0
                                                                       0 eth0
100.1.0.12
                                255.255.255.252 U
                                                                       0 eth3
100.1.4.0
                                255.255.255.0
                                                                       0 eth1
r4:~#
```

 since no routing daemon is currently running, only directly connected destinations are known to the router

starting the routing daemons

on each router (but r5) issue the following command:

```
r4:~# /etc/init.d/zebra start
Starting Zebra daemons (prio:10): zebra ripd.
r4:~#
```

last update: Sept 2018

checking connectivity (again)

towards a remote destination

```
r4:~# ping 100.1.2.1
PING 100.1.2.1 (100.1.2.1) 56(84) bytes of data.
64 bytes from 100.1.2.1: icmp_seq=1 ttl=63 time=0.743 ms
64 bytes from 100.1.2.1: icmp_seq=2 ttl=63 time=0.875 ms
64 bytes from 100.1.2.1: icmp_seq=3 ttl=63 time=0.685 ms

--- 100.1.2.1 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2005ms
rtt min/avg/max/mdev = 0.685/0.767/0.875/0.085 ms
r4:~#
```

after a while, all remote destinations are reachable

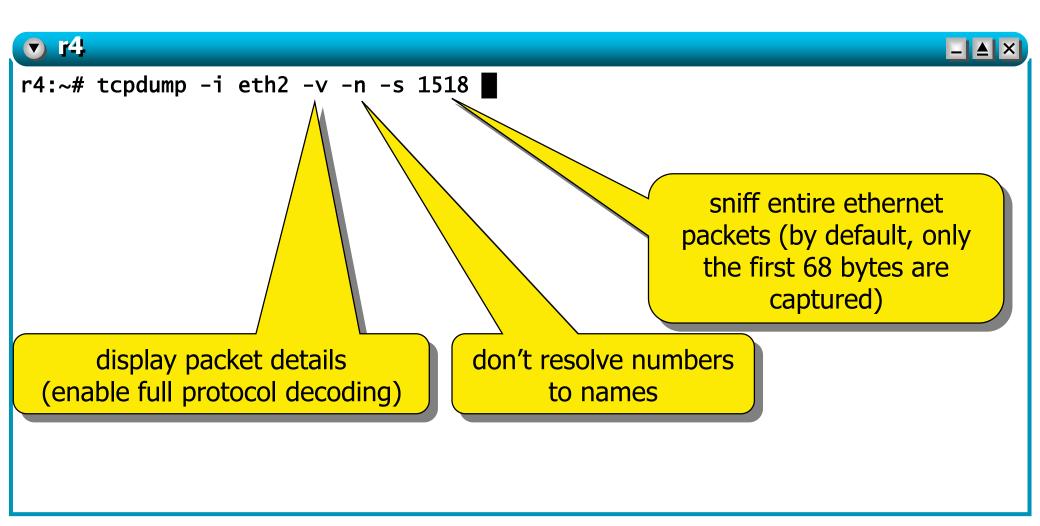
checking the routing table

the routing table is now updated

▼ r4						_ A ×			
r4:~# route									
Kernel IP rout [.]	Kernel IP routing table								
Destination	Gateway	Genmask	Flags	Metric	Ref	Use Iface			
100.1.0.16	*	255.255.255.252	U	0	0	0 eth2			
100.1.0.0	100.1.0.13	255.255.255.252	UG	2	0	0 eth3			
100.1.0.4	100.1.0.17	255.255.255.252	UG	2	0	0 eth2			
100.2.0.0	*	255.255.255.252	U	0	0	0 eth0			
100.1.0.8	100.1.0.17	255.255.255.252	UG	2	0	0 eth2			
100.1.0.12	*	255.255.255.252	U	0	0	0 eth3			
100.1.4.0	*	255.255.255.0	U	0	0	0 eth1			
100.1.2.0	100.1.0.17	255.255.255.0	UG	3	0	0 eth2			
100.1.3.0	100.1.0.17	255.255.255.0	UG	2	0	0 eth2			
100.1.1.0	100.1.0.13	255.255.255.0	UG	2	0	0 eth3			
r4:~#									

a look at ripv2 packets

let's sniff ripv2 packets



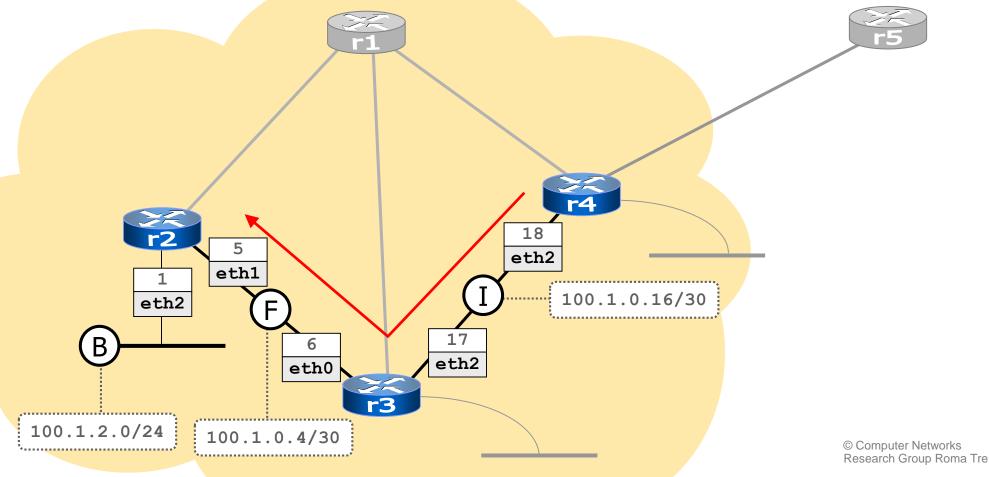
a look at ripv2 packets

let's sniff ripv2 packets

```
▼ r4
r4:~# tcpdump -i eth2 -v -n -s 1518
tcpdump: listening on eth2, link-type EN10MB (Ethernet), capture size 1518
bytes
16:47:48.333986 IP (tos 0x0, ttl 1, id 0, offset 0, flags [DF], length: 152)
100.1.0.17.520 > 224.0.0.9.520: [udp sum ok]
        RIPv2, Response, length: 124, routes: 6
         AFI: IPV4:
                           100.1.0.0/30, tag 0x0000, metric: 2, next-hop: self
                           100.1.0.4/30, tag 0x0000, metric: 1, next-hop: self
         AFI: IPV4:
                           100.1.0.8/30, tag 0x0000, metric: 1, next-hop: self
         AFI: IPV4:
                           100.1.1.0/24, tag 0x0000, metric: 2, next-hop: self
         AFI: IPV4:
                           100.1.2.0/24, tag 0x0000, metric: 2, next-hop: self
         AFI: IPV4:
         AFI: IPV4:
                           100.1.3.0/24, tag 0x0000, metric: 1, next-hop: self
1 packets captured
1 packets received by filter
O packets dropped by kernel
r4:~#
```

a traceroute

```
r4:~# traceroute 100.1.2.1
traceroute to 100.1.2.1 (100.1.2.1), 64 hops max, 40 byte packets
1 100.1.0.17 (100.1.0.17) 10 ms 3 ms 1 ms
2 100.1.2.1 (100.1.2.1) 15 ms 1 ms
r4:~# ■
```



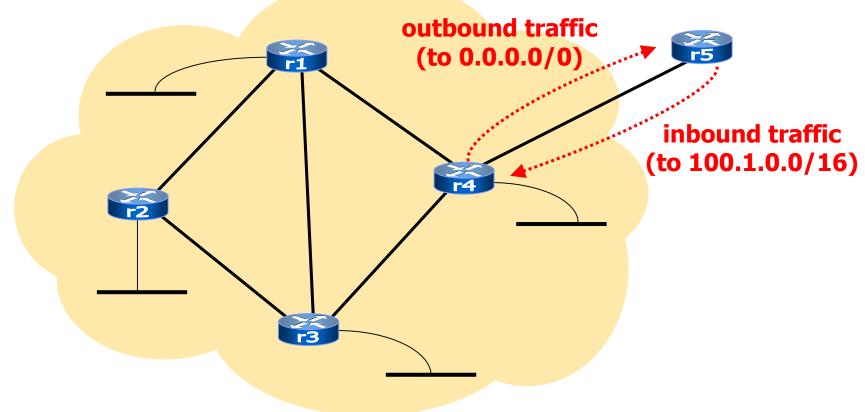
inspecting the rip routing table

```
_ _ ×
r4:~# telnet localhost ripd
Password: zebra
ripd> show ip rip
Codes: R - RIP, C - connected, O - OSPF, B - BGP
      (n) - normal, (s) - static, (d) - default, (r) - redistribute,
      (i) - interface
    Network
                                        Metric From
                                                                Time
                       Next Hop
R(n) 100.1.0.0/30
                        100.1.0.13
                                              2 100.1.0.13
                                                                02:43
R(n) 100.1.0.4/30
                       100.1.0.17
                                                                02:46
                                              2 100.1.0.17
                       100.1.0.17
R(n) 100.1.0.8/30
                                              2 100.1.0.17
                                                                02:46
                                              1 self
C(i) 100.1.0.12/30
                       0.0.0.0
C(i) 100.1.0.16/30
                       0.0.0.0
                                              1 self
R(n) 100.1.1.0/24
                       100.1.0.13
                                              2 100.1.0.13
                                                                02:43
R(n) 100.1.2.0/24
                       100.1.0.17
                                              3 100.1.0.17
                                                                02:46
R(n) 100.1.3.0/24
                       100.1.0.17
                                              2 100.1.0.17
                                                                02:46
C(i) 100.1.4.0/24
                       0.0.0.0
                                              1 self
C(r) 100.2.0.0/30
                                              1 self
                       0.0.0.0
ripd> |
```

© Computer Networks Research Group Roma Tre

static routing

our network is a stub network (i.e., it has just one connection to an external router, r5); hence, static routes are enough for connecting it to the internet



kathara - [lab: rip]

last update: Sept 2018


adding a static route to r5

```
r5:~# route add -net 100.1.0.0/16 gw 100.2.0.1
r5:~# ping 100.1.2.1
PING 100.1.2.1 (100.1.2.1) 56(84) bytes of data.
64 bytes from 100.1.2.1: icmp_seq=1 ttl=62 time=24.1 ms
64 bytes from 100.1.2.1: icmp_seq=2 ttl=62 time=1.11 ms

--- 100.1.2.1 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1023ms
rtt min/avg/max/mdev = 1.117/12.634/24.151/11.517 ms
r5:~#
```

checking connectivity

```
r5:~# traceroute 100.1.2.1
traceroute to 100.1.2.1 (100.1.2.1), 64 hops max, 40 byte packets
1 100.2.0.1 (100.2.0.1) 75 ms 1 ms 2 ms
2 100.1.0.17 (100.1.0.17) 7 ms 1 ms 1 ms
3 100.1.2.1 (100.1.2.1) 24 ms 3 ms 1 ms
r5:~#
```



configuring r4

step 1: configuring the default route

▼ r4.						_ _ ×
r4:~# route add	l default gw 100.	2.0.2				
r4:~# route	_					
Kernel IP routi	ng table					
Destination	Gateway	Genmask	Flags	Metric	Ref	Use Iface
100.1.0.16	*	255.255.255.252	U	0	0	0 eth2
100.1.0.0	100.1.0.13	255.255.255.252	UG	2	0	0 eth3
100.1.0.4	100.1.0.17	255.255.255.252	UG	2	0	0 eth2
100.2.0.0	*	255.255.255.252	U	0	0	0 eth0
100.1.0.8	100.1.0.17	255.255.255.252	UG	2	0	0 eth2
100.1.0.12	*	255.255.255.252	U	0	0	0 eth3
100.1.4.0	*	255.255.255.0	U	0	0	0 eth1
100.1.2.0	100.1.0.17	255.255.255.0	UG	3	0	0 eth2
100.1.3.0	100.1.0.17	255.255.255.0	UG	2	0	0 eth2
100.1.1.0	100.1.0.13	255.255.255.0	UG	2	0	0 eth3
default	100.2.0.2	0.0.0.0	UG	0	0	0 eth0
r4:~#						

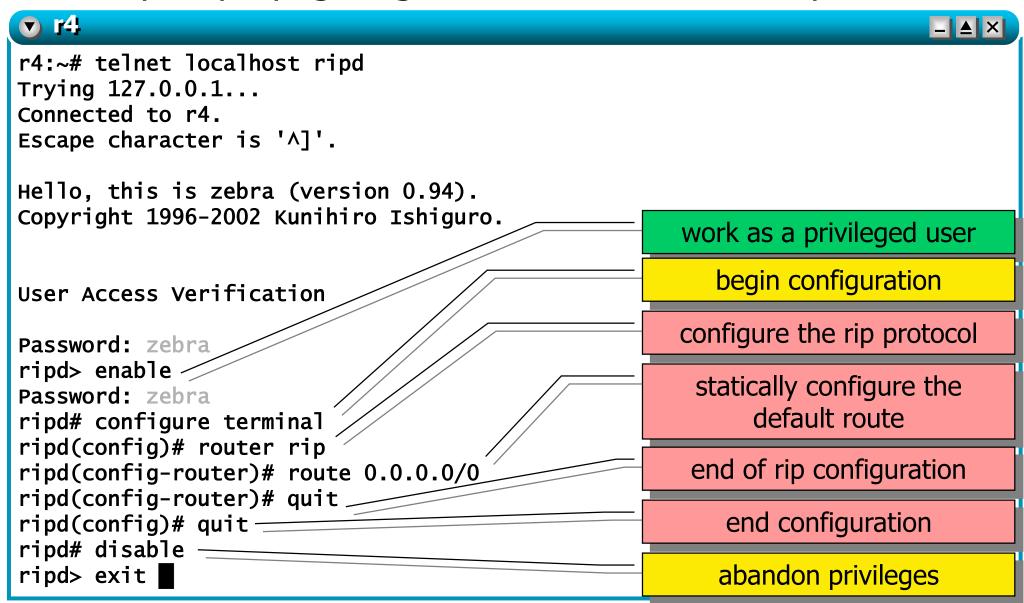
© Computer Networks Research Group Roma Tre

kathara – [lab: rip]

last update: Sept 2018

configuring r4

step 2: propagating the default route into rip



the default route

 after a while, the default route has been injected (via rip) into the network

r1							_ A ×
r1:~# route							
Kernel IP rout	ting table						
Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
100.1.0.16	100.1.0.10	255.255.255.252	UG	2	0	0	eth1
100.1.0.0	*	255.255.255.252	U	0	0	0	eth2
100.2.0.0	100.1.0.14	255.255.255.252	UG	2	0	0	eth0
100.1.0.4	100.1.0.2	255.255.255.252	UG	2	0	0	eth2
100.1.0.8	*	255.255.255.252	U	0	0	0	eth1
100.1.0.12	*	255.255.255.252	U	0	0	0	eth0
100.1.4.0	100.1.0.14	255.255.255.0	UG	2	0	0	eth0
100.1.2.0	100.1.0.2	255.255.255.0	UG	2	0	0	eth2
100.1.3.0	100.1.0.10	255.255.255.0	UG	2	0	0	eth1
100.1.1.0	*	255.255.255.0	U	0	0	0	eth3
default	100.1.0.14	0.0.0.0	UG	2	0	0	eth0
r1:~#							

last update: Sept 2018

checking connectivity

any (even non-existing) destination

r1:~# ping 193.204.161.1

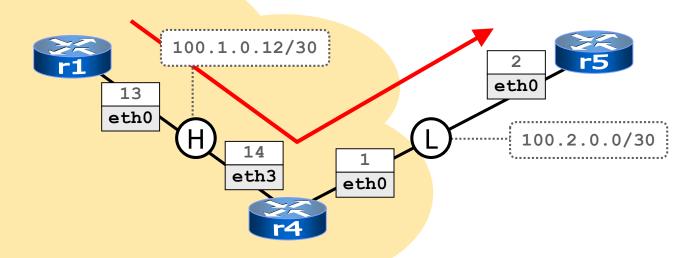
PING 193.204.161.1 (193.204.161.1) 56(84) bytes of data.

From 100.2.0.2 icmp_seq=1 Destination Net Unreachable

From 100.2.0.2 icmp_seq=2 Destination Net Unreachable

--- 193.204.161.1 ping statistics --- 2 packets transmitted, 0 received, +2 errors, 100% packet loss, time 999ms

r1:~#

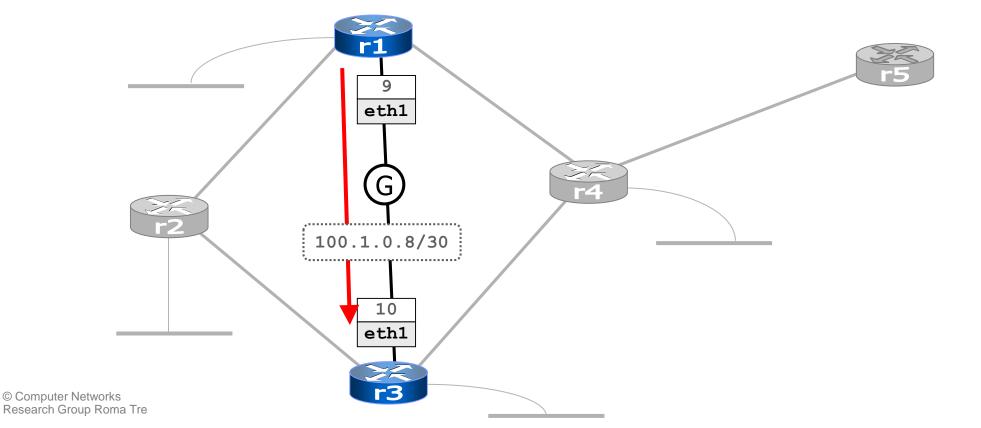


checking connectivity

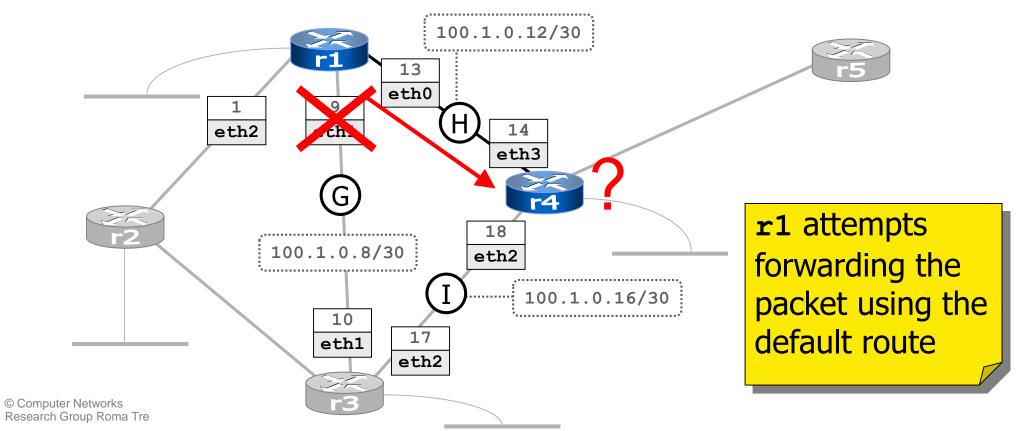
r5 is actually receiving echo request packets

```
▼ r5
                                                                        _ _ ×
r5:~# tcpdump -i eth0 -n -s 1518
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on eth0, link-type EN10MB (Ethernet), capture size 1518 bytes
11:38:43.822503 arp who-has 100.2.0.2 tell 100.2.0.1
11:38:43.824221 arp reply 100.2.0.2 is-at fe:fd:64:02:00:02
11:38:43.825890 IP 100.1.0.13 > 193.204.161.1: icmp 64: echo request seq 1
11:38:43.827139 IP 100.2.0.2 > 100.1.0.13: icmp 92: net 193.204.161.1
unreachable
11:38:44.841566 IP 100.1.0.13 > 193.204.161.1: icmp 64: echo request seq 2
11:38:44.841651 IP 100.2.0.2 > 100.1.0.13: icmp 92: net 193.204.161.1
unreachable
6 packets captured
6 packets received by filter
O packets dropped by kernel
r5:~#
```

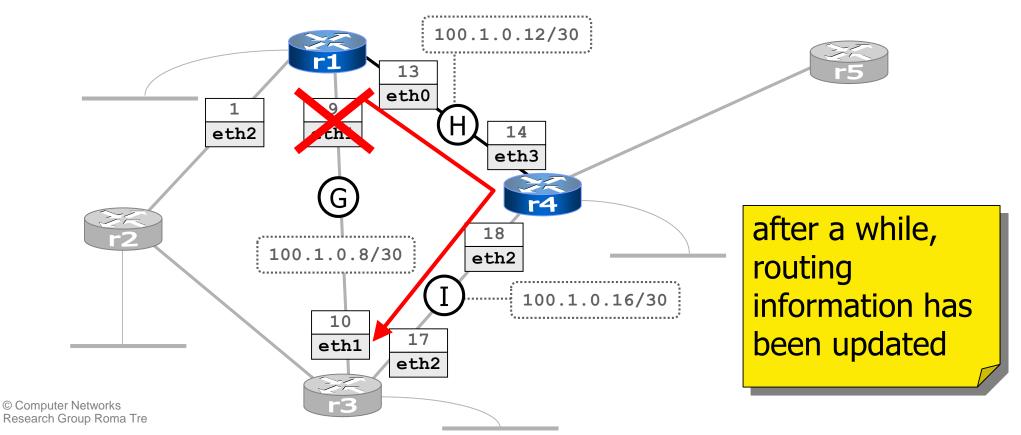
r1:~# traceroute 100.1.0.10 traceroute to 100.1.0.10 (100.1.0.10), 64 hops max, 40 byte packets 1 100.1.0.10 (100.1.0.10) 24 ms 1 ms r1:~# ifconfig eth1 down ■



```
r1:~# traceroute 100.1.0.10
traceroute to 100.1.0.10 (100.1.0.10), 64 hops max, 40 byte packets
1 100.1.0.14 (100.1.0.14) 1 ms 1 ms
2 * * *
3 * * *
```



```
r1:~# traceroute 100.1.0.10
traceroute to 100.1.0.10 (100.1.0.10), 64 hops max, 40 byte packets
1 100.1.0.14 (100.1.0.14) 1 ms 1 ms
2 100.1.0.10 (100.1.0.10) 5 ms 2 ms 1 ms
r1:~#
```



r1's routing table has been updated

▼ r1							- ≜ ×
r1:~# route							
Kernel IP routi	ng table						
Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
100.1.0.16	100.1.0.14	255.255.255.252	UG	2	0	0	eth0
100.1.0.0	*	255.255.255.252	U	0	0	0	eth2
100.2.0.0	100.1.0.14	255.255.255.252	UG	2	0	0	eth0
100-1-0-4	100-1-0-2	255-255-255-252	UG	2	0	0	eth2
100.1.0.8	100.1.0.14	255.255.255.252	UG	3	0	0	eth0
100.1.0.12	×	255.255.255.252	U	U	U	U	eth0
100.1.4.0	100.1.0.14	255.255.255.0	UG	2	0	0	eth0
100.1.2.0	100.1.0.2	255.255.255.0	UG	2	0	0	eth2
100.1.3.0	100.1.0.14	255.255.255.0	UG	3	0	0	eth0
100.1.1.0	*	255.255.255.0	U	0	0	0	eth3
default	100.1.0.14	0.0.0.0	UG	2	0	0	eth0
r1:~#							